Recommended reading

[1] D. Grewe and M. F. P. O'Boyle. *A Static Task Partitioning Approach for Heterogeneous Systems Using OpenCL*. In *CC 2011*, pages 286–305.

[2] K. Kofler, I. Grasso, B. Cosenza, and T. Fahringer. *An Automatic Input-Sensitive Approach for Heterogeneous Task Partitioning*. In *ICS 2013*, pages 149–160, 2013.

[3] C.-K. Luk, S. Hong, and H. Kim. *Qilin: Exploiting Parallelism on Heterogeneous Multiprocessors with Adaptive Mapping*. In *MICRO 2009*, pages 45–55, 2009.

[4] J. Shen, J. Fang, H. Sips, and A. L. Varbanescu. *An application-centric evaluation of OpenCL on multi-core CPUs*. In Parallel Computing, vol. 39, no. 12, pp. 834 – 850, 2013.

[5] Abdullah Gharaibeh, Lauro Beltrão Costa, Elizeu Santos-Neto, and Matei Ripeanu. 2012. *A yoke of oxen and a thousand chickens for heavy lifting graph processing*. In PACT 2012.

[6] Abdullah Gharaibeh, Lauro Beltrao Costa, Elizeu Santos-Neto, and Matei Ripeanu. 2013. *On Graphs, GPUs, and Blind Dating: A Workload to Processor Matchmaking Quest*. In IPDPS 2013.

[7] Thomas Scogland, Barry Rountree, Wu-chun Feng, Bronis R. de Supinski: *Heterogeneous Task Scheduling for Accelerated OpenMP*. In IPDPS 2012.

[8] M. Viñas, B.B. Fraguela, Z. Bozkus, D. Andrade. *Improving OpenCL programmability with the Heterogeneous Programming Library*. In ICCS 2015.

[9] Pieter Hijma, Ceriel Jacobs, Rob van Nieuwpoort, and Henri Bal: *Cashmere: Heterogeneous Many-Core Computing*. In IPDPS 2015.

[10] Jie Shen and Ana Lucia Varbanescu and Yutong Lu and Peng Zou and Henk Sips (2016). *Workload Partitioning for Accelerating Applications on Heterogeneous Platforms*. IEEE Transactions on Parallel and Distributed Systems.

[11] Jie Shen and Ana Lucia Varbanescu and Xavier Martorell and Henk Sips (2015). *A Study of Application Kernel Structure for Data Parallel Applications*. Technical Report PDS-2015-001, Delft University of Technology.

[12] Jie Shen and Ana Lucia Varbanescu and Xavier Martorell and Henk Sips (2015). *Matchmaking Applications and Partitioning Strategies for Efficient Execution on Heterogeneous Platforms*. In ICPP 2015.

[13] Jie Shen and Ana Lucia Varbanescu and Henk Sips (2014). *Look Before You Leap: Using the Right Hardware Resources to Accelerate Applications*. In HPCC 2014.

[14] Jie Shen and Ana Lucia Varbanescu and Peng Zou and Yutong Lu and Henk Sips (2014). *Improving Performance by Matching Imbalanced Workloads with Heterogeneous Platforms*. In ICS 2014.

[15] Yong Guo, Ana Lucia Varbanescu, Dick Epema and Alexandru Iosup. *Design and Experimental Evaluation of Distributed Heterogeneous Graph-Processing Systems*. In CCGrid 2016.

[16] Jie Shen and Ana Lucia Varbanescu and Henk Sips and Michael Arntzen and Dick Simons (2013). *Glinda: A Framework for Accelerating Imbalanced Applications on Heterogeneous Platforms*. In CF 2013.

[17] Souley Madougou, Ana Varbanescu, Cees de Laat, Rob van Nieuwpoort (2016). *The Landscape of GPGPU Performance Modeling Tools*. Journal of Parallel Computing, 2016.

[18] Pieter Hijma, Rob V. van Nieuwpoort, Ceriel J.H. Jacobs, and Henri E.Bal. *Stepwise-refinement for performance: a methodology for many-core programming.* Concurrency and Computation: Practice and Experience, 2015.

[19] Rob V. van Nieuwpoort, Gosia Wrzesinska, Ceriel J. H. Jacobs, and Henri E. Bal. Satin: A High-Level and Efficient Grid Programming Model. ACM TOPLAS, 2010.

[20] M. Viñas, Z. Bozkus, B.B. Fraguela. *Exploiting heterogeneous parallelism with the Heterogeneous Programming Library*. JPDC, 73(12):1627-1638. 2013

[21] J.F. Fabeiro, D. Andrade, B.B. Fraguela. *Writing a performance-portable matrix multiplication*. Parallel Computing, 52:65-77. 2016

[22] M. Viñas, B.B. Fraguela, D. Andrade, R. Doallo. *High Productivity Multi-device Exploitation with the Heterogeneous Programming Library*. J. Parallel and Distributed Computing, 101:51-68. 2017

[23] Stijn Heldens, Ana Lucia Varbanescu and Alexandru Iosup. *Dynamic Load Balancing for High-Performance Graph Processing on Hybrid CPU-GPU Platforms*. SC-Workshop'16.